Ayesha Raza

AI/ML Engineer - LLM apps, RAG and production ML

Lahore, Pakistan · 6 years of freelance experience

Skills

AWS SageMaker, Airflow, Celery, Docker, FAISS, FastAPI, Hugging Face Transformers, LLM fine-tuning, LangChain, LangGraph, MLOps, MLflow, NLP, ONNX, OpenAI API, OpenCV, PyTorch, Python, RAG, XGBoost, YOLO, computer vision, pandas, pgvector, prompt engineering, recommendation systems, scikit-learn, spaCy, time series forecasting

Selected projects

• RAG chatbot over logistics documents for a logistics client (2020): FAISS, Python; served 38k users in the first month.
• Multi-agent workflow for real estate operations for a real estate client (2021): OpenAI API, prompt engineering; reduced cloud cost by 30%.
• Fine-tune an open LLM for real estate support tickets for a real estate client (2023): PyTorch, LLM fine-tuning; raised conversion by 50%.
• Churn / risk prediction model for a insurance company for a insurance client (2024): XGBoost, Docker; cut manual review time by 47%.
• Demand forecasting for logistics for a logistics client (2025): Python, time series forecasting; raised conversion by 19%.
• Defect detection with computer vision for healthcare for a healthcare client (2026): ONNX, Python; raised conversion by 43%.
• Document image classification for e-commerce for a e-commerce client (2025): FastAPI, Python; brought latency down to 80 ms p95.
• Entity extraction from logistics contracts for a logistics client (2020): Python, FastAPI; saved the team 17 hours per week.
• Productionise ML models for a e-commerce team for a e-commerce client (2020): MLflow, AWS SageMaker; improved accuracy from 83% to 89%.
• Recommendation engine for a manufacturing platform for a manufacturing client (2020): FastAPI, recommendation systems; brought latency down to 80 ms p95.
• Productionise ML models for a media team for a media client (2020): Celery, Docker; served 40k users in the first month.
• Entity extraction from healthcare contracts for a healthcare client (2023): Python, spaCy; shipped in 10 weeks.
• Productionise ML models for a fintech team for a fintech client (2026): Celery, Docker; brought latency down to 250 ms p95.
• Defect detection with computer vision for real estate for a real estate client (2023): computer vision, YOLO; improved accuracy from 72% to 79%.